Real science training resource comparison
CPU task containers: this run does not measure GPU-container packing or allocator fragmentation.

Completed updates vs measured step time

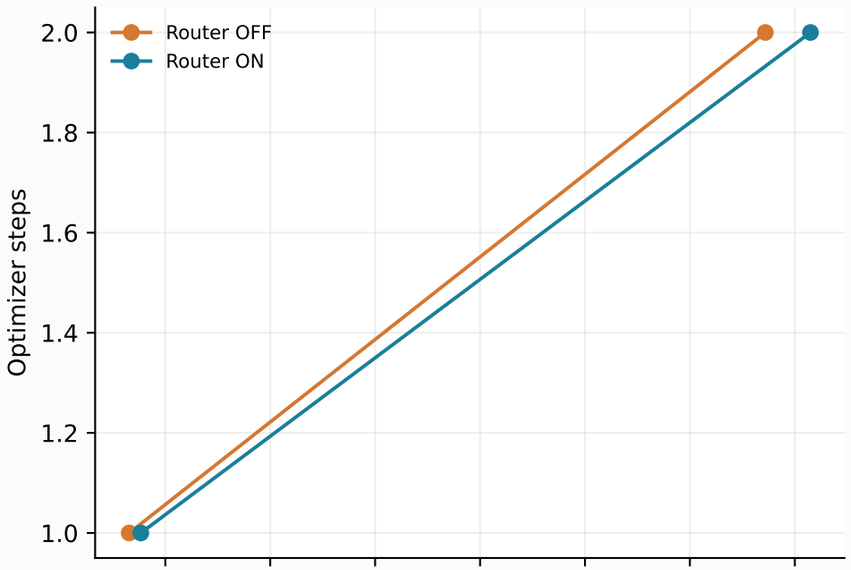

GPU utilization: mean of 3 GPUs

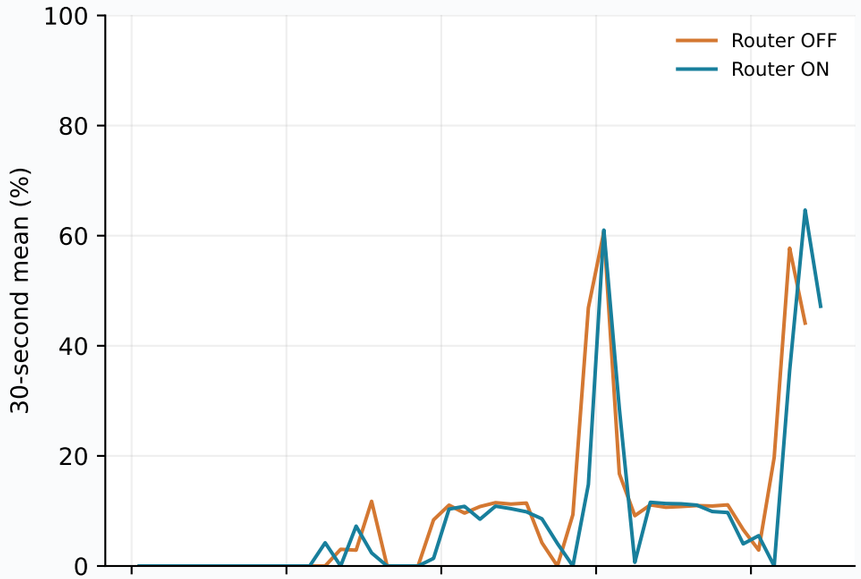

Device memory used: sum of 3 GPUs

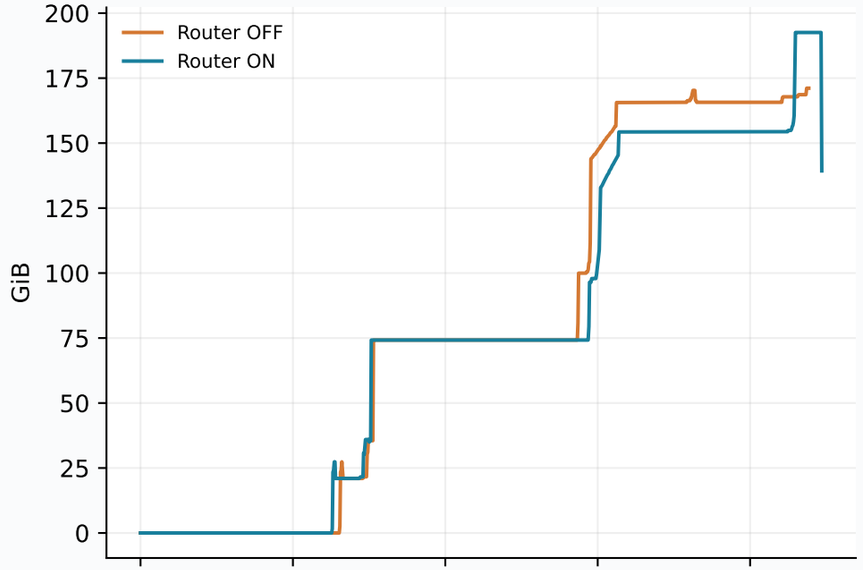

Training allocator peaks (not fragmentation)

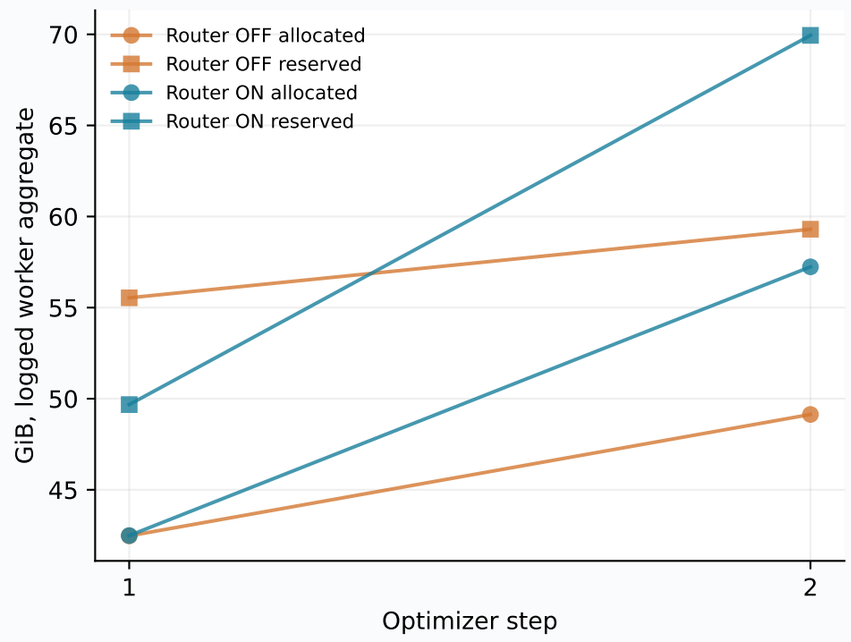

Router ON admission accounting

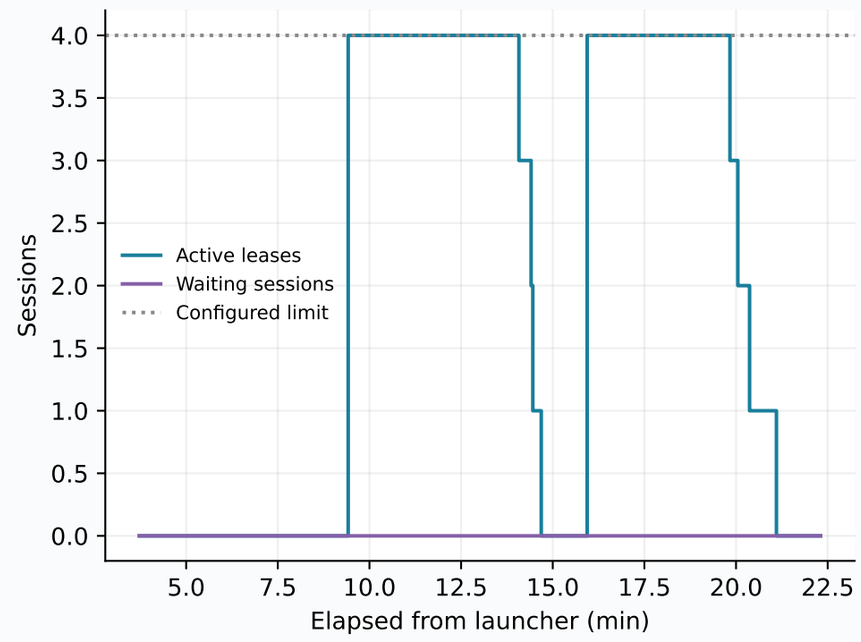
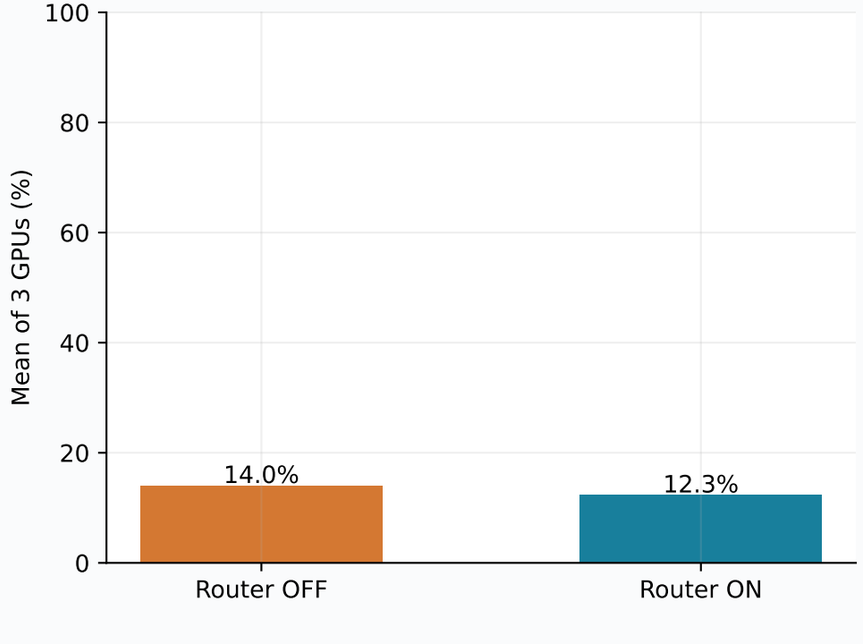
Mean utilization while GPU memory > 1 GiB