Real LAPS repair training | Qwen3.5-4B | identical task / model / concurrency budgets
2 requested steps per arm. Short-run loss and entropy do not establish learning superiority.

Policy entropy

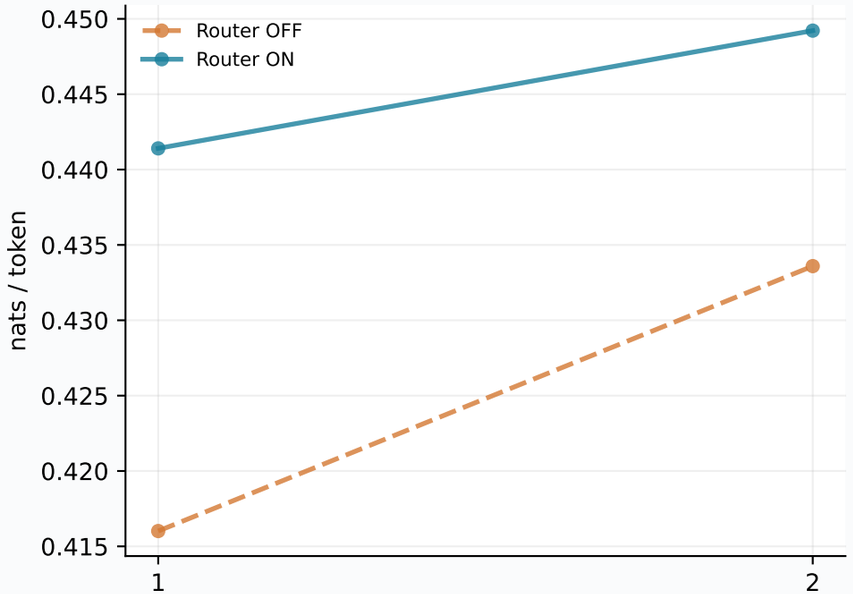

Policy gradient loss

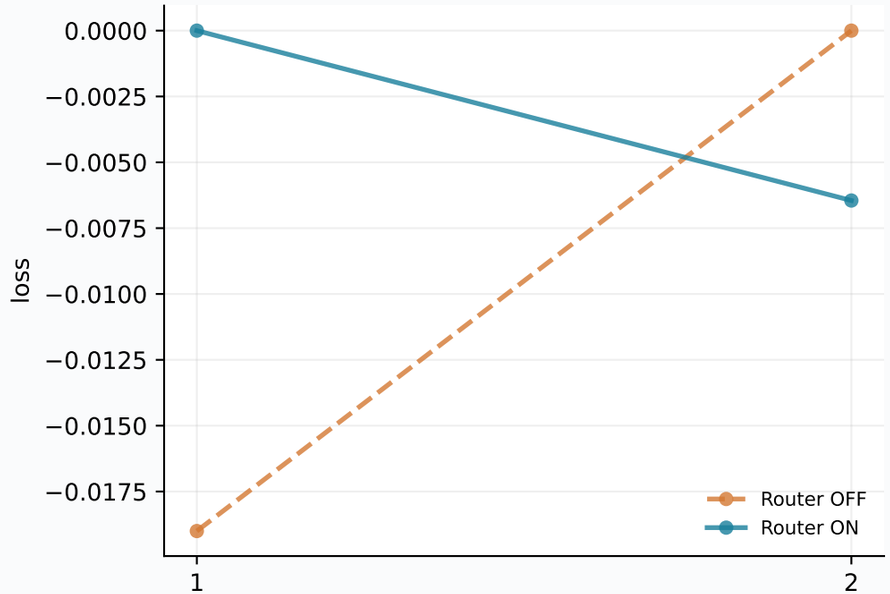

Gradient norm

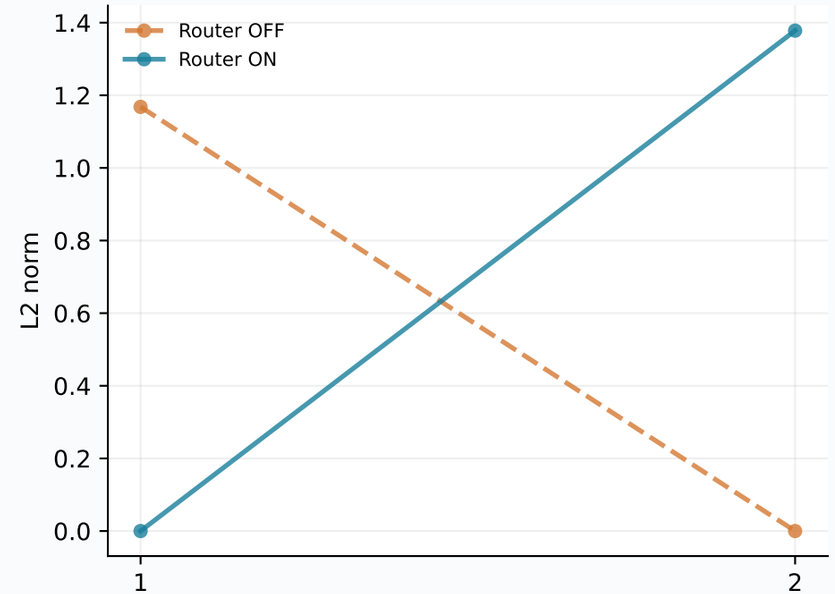

Task reward

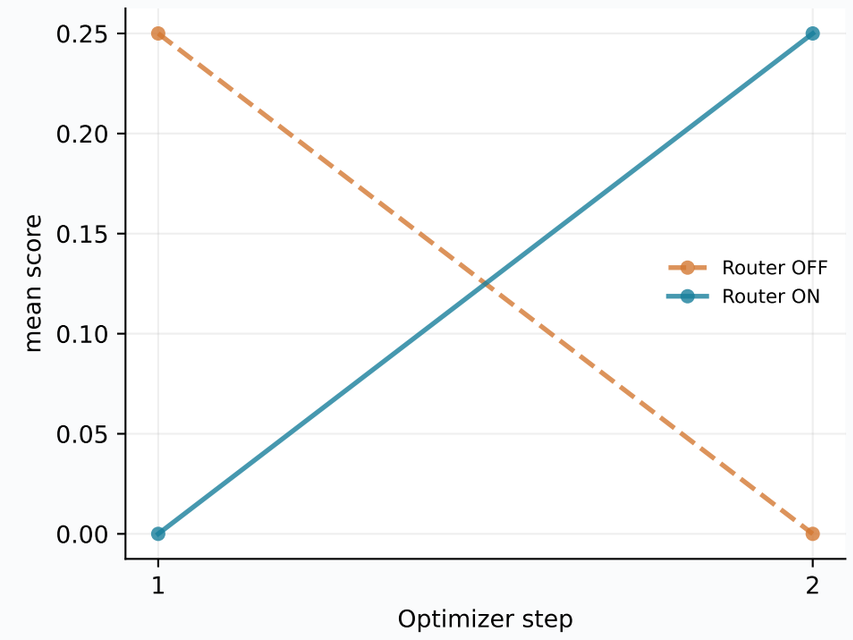

Masked training samples

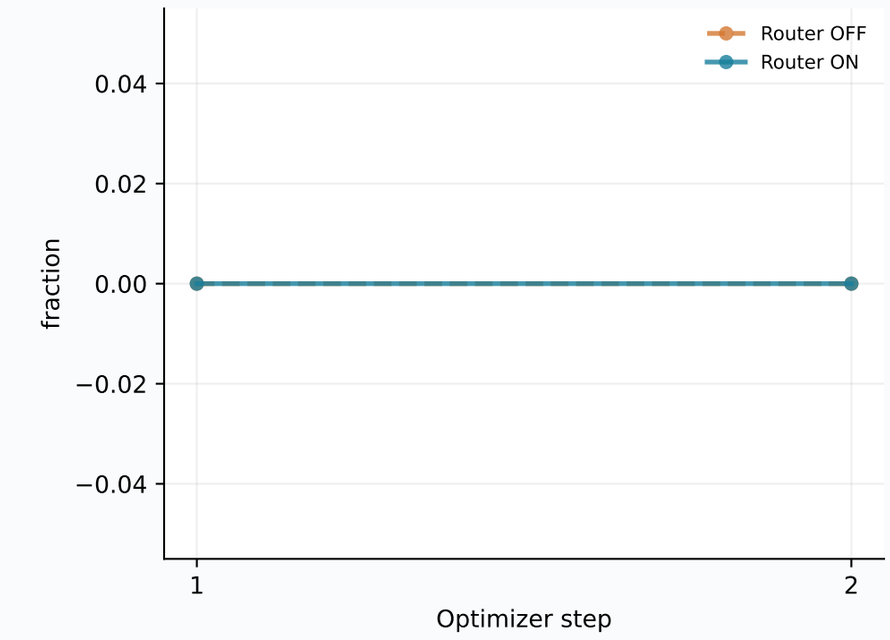
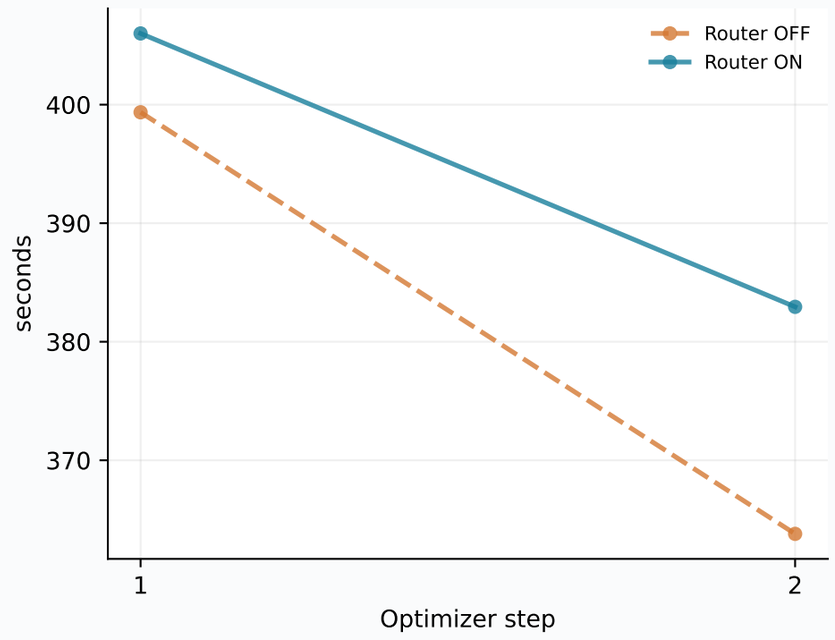
Time per training step